


MHC Transformer Block

Single block with 3-channel residual dynamics
 x_0 skip + bigram inject



Key: The sublayer (attention or MLP) sees the *full* d_{model} -dim input.
 Channel separation is *only* on the residual update side via learned γ_i gains.
 Compute channel ($\gamma \approx 1$) gets full updates. Memory channel ($\gamma \approx 0.1$) changes slowly.